


Toward Controlled Generation of Text

Hu, Yang, Liang, Salakhutdinov & Xing

ICML 2017

Motivation



1

¹https://commons.wikimedia.org/wiki/File:Autoencoder_structure.png

Distentangled Latent Representations

I love this movie →

I hate this movie →

0.21	0.32	0.74	0.43
0.45	0.78	0.97	0.17

Distentangled Latent Representations

I love this movie →

I hate this movie →

0.21	0.32	0.74	0.43
0.45	0.78	0.97	0.17



I love this movie →

I hate this movie →

0.68	0.12	0.33	1.00
0.68	0.12	0.33	0.00

Problem Statement

Generate fake samples similar to the source distribution by conditioning their generation on a tunable set of attributes.

Notation

$x \Rightarrow$ **source corpus**

$\hat{x} \Rightarrow$ **output corpus**

$c \Rightarrow$ **structured code**, known label for each document

$z \Rightarrow$ **unstructured latent code**

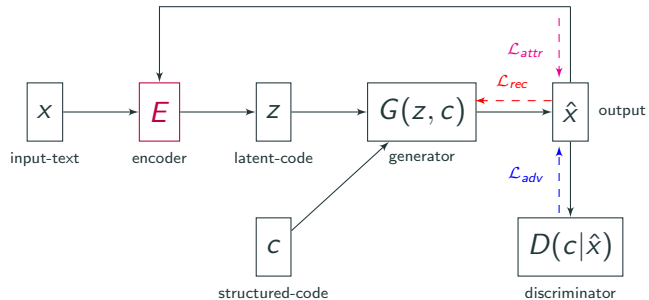
$E \Rightarrow$ **encoder**, parameterized to generate z

$G \Rightarrow$ **decoder/generator**, produces \hat{x} conditioned on (z, c)


$D \Rightarrow$ **discriminator**, predicts c given \hat{x} .

$\tau \Rightarrow$ **softmax temperature**, for decoder word prediction

Architecture



Unrolled Architecture



Discriminator Optimization

- Maximize the likelihood of predicting the correct distribution of the structured code c given the labeled examples X_L .

$$\mathcal{L}_s(\theta_D) = -\mathbb{E}_{X_L}[\log q_D(c_L|x_L)]$$

- Maximize the likelihood of predicting the correct distribution of the structured code c given the generated sentences \hat{x} . Also minimize the empirically observed Shannon entropy of the discriminator predictions $q_D(c'|\hat{x})$.

$$\mathcal{L}_u(\theta_D) = -\mathbb{E}_{p_G(\hat{x}|z,c)p(z)p(c)}[\log q_D(c|\hat{x}) + \beta \mathcal{H}(q_D(c'|\hat{x}))]$$

Generator Optimization

- Maximize the likelihood of predicting the original document x , given the latent spaces and the generator $G(z, c)$.

$$\begin{aligned}\mathcal{L}_{VAE}(\theta_G, \theta_E; x) = & -\mathbb{E}_{q_E(z|x)q_D(c|x)}[\log p_G(x|z, c)] \\ & + KL(q_E(z|x)||p(z))\end{aligned}$$

- Maximize the likelihood of generating the output documents with the correct structured code c

$$\mathcal{L}_{attr,c}(\theta_G) = -\mathbb{E}_{p(z)p(c)}[\log q_D(c|\tilde{G}_\tau(z, c))]$$

Additional Generator Optimization: Independency Constraint

The encoder is re-used to regenerate the latent distribution z devoid of the structured code c , from the output distribution $\tilde{G}_\tau(z, c)$.

$$\mathcal{L}_{attr,z}(\theta_G) = -\mathbb{E}_{p(z)p(c)}[\log q_E(z|\tilde{G}_\tau(z, c))]$$

Training Objectives

Generator:

$$\begin{aligned}\min_{\theta_G} \mathcal{L}_G = & \quad \mathcal{L}_{VAE} \quad (\text{reconstruction loss}) \\ & + \lambda_c \mathcal{L}_{attr,c} \quad (\text{style entanglement}) \\ & + \lambda_z \mathcal{L}_{attr,z} \quad (\text{independency constraint})\end{aligned}$$

Discriminator:

$$\begin{aligned}\min_{\theta_D} \mathcal{L}_D = & \quad \mathcal{L}_s \quad (\text{labeled example classification}) \\ & + \lambda_u \mathcal{L}_{attr,u} \quad (\text{synthesized example classification})\end{aligned}$$

Require: A large corpus of unlabeled sentences $\mathcal{X} = \{x\}$


A few sentence attribute labels $\mathcal{X}_L = \{(x_L, c_L)\}$

Parameters: $\lambda_c, \lambda_z, \lambda_u, \beta$ – balancing parameters

- 1: Initialize the base VAE by minimizing \mathcal{L}_{VAE} on \mathcal{X}
- 2: **repeat**
- 3: Train the discriminator D by minimizing \mathcal{L}_u
- 4: Train the generator G and the encoder E by minimizing \mathcal{L}_G and \mathcal{L}_{VAE} , respectively.
- 5: **until** convergence

Ensure: Sentence generator G conditioned on disentangled representation (z, c)

Wake-Sleep Algorithm



Experimental Setup

Simple Reconstruction

- 350K IMDB movie reviews
- Sentence length ≤ 15
- Total sentence count = 1.4M
- Vocab size = 16K

Binary Sentiment Classification and Conditioned Generation

- **Stanford Sentiment Treebank-2:** Movie reviews
 - 2837 training examples
 - Sentence length ≤ 15
- **Sentiment Lexicon:** Words used as sentences
 - 2700 words
 - Sentence length ≤ 15
- **IMDB:** Movie reviews
 - 16K training examples

Tense Classification and Conditioned Generation

- TimeBank² Lexicon
- 5250 words and phrases labeled with one of past, present, future
- Verbs in different tenses (e.g., was, will be) as well as time expressions (e.g., in the future)


²<http://timeml.org>

Parameters

- The generator and encoder are set as single-layer LSTM RNNs with input/hidden dimension of 300 and max decoding time-step count of 15
- Discriminators are set as Convolutional Nets
- VAE KL-Divergence weight linearly annealed from 0 to 1 during training
- Softmax temperature τ annealed from 1 to 0
- Balancing λ weights all set to 0.1, β is selected on the dev set

Results

Results



Without 'Independency Constraint'

the acting is bad

⇒

the movie is so much fun

none of this is very original

⇒

highly recommended viewing for its
courage , and ideas

too bland

⇒

highly watchable

i can analyze this movie without more
than three words

⇒

i highly recommend this film to anyone
who appreciates music

With 'Independency Constraint'

the film is strictly routine	⇒	the film is full of imagination
after watching this movie , i felt that disappointed	⇒	after seeing this film , i 'm a fan
the acting is uniformly bad either	⇒	the performances are uniformly good
this is just awful	⇒	this is pure genius

Varying the unstructured code z

(“negative”, “past”)

the acting was also kind of hit or miss .

(“positive”, “past”)

his acting was impeccable

(“negative”, “present”)

the era seems impossibly distant

(“positive”, “present”)

i 've always been a big fan of the smart dialogue .

(“negative”, “future”)

and that would be devastating !

(“positive”, “future”)

i will definitely be buying this on dvd

-
- Non-parallel attribute-controlled generation
 - Encoder independency constraint

-
- No quantifiable results for content preservation

Related Work

- Shen, Tianxiao, et al. 'Style transfer from non-parallel text by cross-alignment.' Advances in Neural Information Processing Systems. 2017.
- Kim, Yoon, et al. 'Adversarially regularized autoencoders for generating discrete structures.' arXiv preprint arXiv:1706.04223 (2017).
- Fu, Zhenxin, et al. 'Style Transfer in Text: Exploration and Evaluation.' arXiv preprint arXiv:1711.06861 (2017).
- Melnyk, Igor, et al. 'Improved Neural Text Attribute Transfer with Non-parallel Data.' arXiv preprint arXiv:1711.09395 (2017).

Questions?